

# All Codes:

## TASK 1:

### Mapper:

```
#!/usr/bin/env python
```

```
import sys
```

```
# Initialize a dictionary to store trip count and total distance per taxi
```

```
taxi_data = {}
```

```
for line in sys.stdin:
```

```
    # Split the line into fields
```

```
    fields = line.strip().split(',')
```

```
    # Check if the line has the expected number of fields
```

```
    if len(fields) >= 8:
```

```
        try:
```

```
            # Parse trip information
```

```
            taxi_id = fields[1]
```

```
            distance = float(fields[3])
```

```
        # Update the dictionary with trip count and total distance
```

```
        if taxi_id in taxi_data:
```

```
            trip_count, total_distance = taxi_data[taxi_id]
```

```
            trip_count += 1
```

```
            total_distance += distance
```

```
            taxi_data[taxi_id] = (trip_count, total_distance)
```

```
        else:
```

```
            taxi_data[taxi_id] = (1, distance)
```

```
except ValueError:
    # Skip lines with invalid data
    continue
```

```
# Emit intermediate key-value pairs from the mapper
for taxi_id, (trip_count, total_distance) in taxi_data.items():
    print(f"{taxi_id}\t{trip_count},{total_distance}")
```

### **Reducer:**

```
#!/usr/bin/env python
```

```
import sys
```

```
# Initialize variables to store trip count and total distance per taxi
```

```
current_taxi = None
```

```
trip_count = 0
```

```
total_distance = 0
```

```
for line in sys.stdin:
```

```
    # Split the line into key and value
```

```
    key, value = line.strip().split('\t')
```

```
    taxi_id, trip_info = key, value
```

```
    # Split trip_info into trip count and total distance
```

```
    count, distance = map(float, trip_info.split(','))
```

```
    if current_taxi == taxi_id:
```

```
        # Accumulate trip count and total distance for the current taxi
```

```
        trip_count += count
```

```
        total_distance += distance
```

else:

if current\_taxi:

# Calculate and emit the average distance per trip for the previous taxi

average\_distance = total\_distance / trip\_count

print(f"{current\_taxi}\t{trip\_count},{average\_distance}")

# Reset variables for the new taxi

current\_taxi = taxi\_id

trip\_count = count

total\_distance = distance

# Output the result for the last taxi

if current\_taxi:

average\_distance = total\_distance / trip\_count

print(f"{current\_taxi}\t{trip\_count},{average\_distance}")

### Shell Script:

#!/bin/bash

# Hadoop Streaming jar location

HADOOP\_STREAMING\_JAR="/usr/lib/hadoop/hadoop-streaming-2.10.1-amzn-1.1.jar"

# Input and output paths

INPUT\_PATH="/input/Trips.txt"

OUTPUT\_PATH="/output/task1"

# Remove previous output

hadoop fs -rm -r \$OUTPUT\_PATH

# Run Hadoop Streaming job

```
hadoop jar $HADOOP_STREAMING_JAR \  
-D stream.map.output.field.separator="\t" \  
-D stream.num.map.output.key.fields=2 \  
-D mapreduce.job.reduces=3 \  
-files task1_mapper.py,task1_reducer.py \  
-mapper "python3 task1_mapper.py" \  
-reducer "python3 task1_reducer.py" \  
-input $INPUT_PATH \  
-output $OUTPUT_PATH
```

# Display the output

```
hadoop fs -cat $OUTPUT_PATH/part-*
```

## **TASK 2:**

### **Mapper:**

```
#!/usr/bin/env python
```

```
import sys
```

```
import random
```

```
# Read k and v from command line arguments
```

```
k = int(sys.argv[1])
```

```
v = int(sys.argv[2])
```

```
# Initialize medoids with random data points
```

```
medoids = []
```

```
data = []
```

```
for line in sys.stdin:
```

```
    fields = line.strip().split(',')
```

```
trip_id = int(fields[0])
pickup_x = float(fields[4])
pickup_y = float(fields[5])

if len(medoids) < k:
    medoids.append((trip_id, pickup_x, pickup_y))
data.append((trip_id, pickup_x, pickup_y))
```

# Assignment step

for point in data:

```
trip_id, pickup_x, pickup_y = point
min_distance = float("inf")
closest_medoid = None
```

for medoid in medoids:

```
medoid_id, medoid_x, medoid_y = medoid
distance = (pickup_x - medoid_x) ** 2 + (pickup_y - medoid_y) ** 2
```

```
if distance < min_distance:
    min_distance = distance
    closest_medoid = medoid_id
```

```
print(f"{closest_medoid}\t{trip_id},{pickup_x},{pickup_y}")
```

# Output data for reducer

for point in data:

```
trip_id, pickup_x, pickup_y = point
print(f"{trip_id}\t{pickup_x},{pickup_y}")
```

**Reducer:**

```
#!/usr/bin/env python
```

```
import sys
```

```
# Identity reducer
```

```
for line in sys.stdin:
```

```
    print(line.strip())
```

**Shell Script:**

```
#!/bin/bash
```

```
# Set input and output paths
```

```
input_path="/input/Trips.txt"
```

```
output_path="/output/task2"
```

```
# Hadoop Streaming jar location
```

```
HADOOP_STREAMING_JAR="/usr/lib/hadoop/hadoop-streaming-2.10.1-amzn-1.1.jar"
```

```
# Remove previous output (if any)
```

```
hadoop fs -rm -r $output_path
```

```
# Run the MapReduce job
```

```
hadoop jar $HADOOP_STREAMING_JAR \
```

```
    -D mapreduce.job.reduces=3 \
```

```
    -files task2_mapper.py,task2_reducer.py \
```

```
    -mapper "python3 task2_mapper.py $1 $2" \
```

```
    -reducer "python3 task2_reducer.py" \
```

```
    -input $input_path \
```

```
    -output $output_path
```

```
# Display the final results  
hadoop fs -cat $output_path/part-*
```

### **TASK 3:**

#### **Mapper 1:**

```
#!/usr/bin/env python  
import sys  
  
# Loop through input lines  
for line in sys.stdin:  
    line = line.strip()  
    data = line.split(',')  
  
    # Check if the data has 4 elements and the first element is a digit  
    if len(data) == 4 and data[0].isdigit():  
        taxi_id, company, model, year = data  
        print(f"{taxi_id}\tT\t{company}")  
  
# Loop through input lines  
for line in sys.stdin:  
    line = line.strip()  
    data = line.split(',')  
  
    # Check if the data has 8 elements and the first element is a digit  
    if len(data) == 8 and data[0].isdigit():  
        trip_id, taxi_id, fare, distance, pickup_x, pickup_y, dropoff_x, dropoff_y = data  
        print(f"{taxi_id}\tR\t1")
```

**Reducer 1:**

```
#!/usr/bin/env python

import sys

current_taxi_id = None
current_company = None
trip_count = 0

# Loop through input lines
for line in sys.stdin:
    line = line.strip()
    taxi_id, data_type, value = line.split('\t')

    # Check if the current taxi_id is different from the previous one
    if current_taxi_id != taxi_id:
        if current_taxi_id:
            # Emit data in the format: taxi_id current_company
            print(f"{current_taxi_id}\t{current_company}")

        current_taxi_id = taxi_id
        current_company = None
        trip_count = 0

    # Determine if it's a 'T' (Taxi) or 'R' (Record) data type
    if data_type == 'T':
        current_company = value
    elif data_type == 'R':
        trip_count += int(value)

# Print the last taxi's data
if current_taxi_id:
```



```
print(f'{current_taxi_id}\t{current_company}')
```

### **Mapper 2:**

```
#!/usr/bin/env python
```

```
import sys
```

```
# Loop through input lines
```

```
for line in sys.stdin:
```

```
    line = line.strip()
```

```
# Split the input line by the tab character ('\t')
```

```
taxi_id, company = line.split('\t')
```

```
# Emit data in the format: company 1
```

```
print(f'{company}\t1')
```

### **Reducer 2:**

```
#!/usr/bin/env python
```

```
import sys
```

```
current_company = None
```

```
trip_count = 0
```

```
# Loop through input lines
```

```
for line in sys.stdin:
```

```
    line = line.strip()
```

```
    company, count = line.split('\t')
```

```
# Check if the current company is different from the previous one
```

```

if current_company != company:

    if current_company:

        # Emit data in the format: current_company    trip_count
        print(f"{current_company}\t{trip_count}")

    current_company = company

    trip_count = 0

# Accumulate trip counts
trip_count += int(count)

# Print the last company's trip count
if current_company:
    print(f"{current_company}\t{trip_count}")

```

### **Shell Script:**

```

#!/bin/bash

# Hadoop Streaming jar location
HADOOP_STREAMING_JAR="/usr/lib/hadoop/hadoop-streaming-2.10.1-amzn-1.1.jar"

# Define input and output directories
INPUT_PATH="/input"
OUTPUT_PATH="/output/task3"

# Remove previous output (if any)
hadoop fs -rm -r $OUTPUT_PATH
hadoop fs -rm -r /intermediate_output

# Run the first MapReduce job for the join operation

```

```
hadoop jar $HADOOP_STREAMING_JAR \  
-D stream.map.output.field.separator="\t" \  
-D stream.num.map.output.key.fields=2 \  
-D mapreduce.job.reduces=3 \  
-files task3_mapper1.py,task3_reducer1.py \  
-mapper "python3 task3_mapper1.py" \  
-reducer "python3 task3_reducer1.py" \  
-input $INPUT_PATH/Trips.txt $INPUT_PATH/Taxis.txt \  
-output /intermediate_output
```

# Run the second MapReduce job for counting

```
hadoop jar $HADOOP_STREAMING_JAR \  
-D stream.map.output.field.separator="\t" \  
-D stream.num.map.output.key.fields=2 \  
-D mapreduce.job.reduces=3 \  
-files task3_mapper2.py,task3_reducer2.py \  
-mapper "python3 task3_mapper2.py" \  
-reducer "python3 task3_reducer2.py" \  
-input /intermediate_output/part-* \  
-output $OUTPUT_PATH
```

# Display the final results

```
hadoop fs -cat $OUTPUT_PATH/part-*
```